

## ▼ Download Data

```

1  !wget http://www2.aueb.gr/users/ion/data/lingspam_public.tar.gz
2  !tar -zxf lingspam_public.tar.gz

--2021-11-01 19:27:09--  http://www2.aueb.gr/users/ion/data/lingspam_public.tar.gz
Resolving www2.aueb.gr (www2.aueb.gr)... 195.251.255.138
Connecting to www2.aueb.gr (www2.aueb.gr)|195.251.255.138|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 11564714 (11M) [application/x-gzip]
Saving to: 'lingspam_public.tar.gz'

lingspam_public.tar 100%[=====>] 11.03M  1.73MB/s   in 7.0s

2021-11-01 19:27:17 (1.59 MB/s) - 'lingspam_public.tar.gz' saved [11564714/11564714]

```

## ▼ Feature selection using information gain(IG) matrix

```

1  import os
2  import re
3  import math
4  import numpy as np
5
6  from collections import Counter
7
8  # Copy from readme.txt
9  # Each one of the 10 subdirectories contains both spam and legitimate
10 # messages, one message in each file. Files whose names have the form
11 # spmsg*.txt are spam messages. All other files are legitimate messages.
12
13 path = "lingspam_public/lemm_stop"
14 word_filter = re.compile(r'^[a-z]+--[a-z]+[0-9]*$')
15
16 spam_words = []
17 ham_words = []
18 train_X = []
19 train_Y = []
20 spam_num = 0
21 ham_num = 0
22
23 # Build IG and training dataset
24 for dir in [os.path.join(path, 'part' + str(i)) for i in range(1, 10)]:
25     mails = [(os.path.join(dir, file_name), 'spmsg' in file_name) for file_name in os.li
26
27     for file_path, spam in mails:

```

```

28         with open(file_path) as f:
29             content = f.readlines()[2]
30             words = content.split()
31             filtered_words = list(set([w for w in words if (len(w) > 1 and re.match(word_filters, w))]))
32             # Add words to training set
33             train_X.append(filtered_words)
34             # Label spam as 1, ham as 0
35             if spam:
36                 train_Y.append(1)
37                 spam_words += filtered_words
38                 spam_num += 1
39             else:
40                 train_Y.append(0)
41                 ham_words += filtered_words
42                 ham_num += 1
43
44
45 train_all_counter = Counter(spam_words + ham_words)
46 train_spam_counter = Counter(spam_words)
47 train_ham_counter = Counter(ham_words)
48
49 print(f'Loading {spam_num + ham_num} emails, {spam_num} of them are spam, \
50 and {ham_num} of them are ham.')
51 print(f'Within the total of {len(train_all_counter)} words, \
52 getting {len(train_spam_counter)} words occur in spam, \
53 and {len(train_ham_counter)} words occur in ham email.')

```

Loading 2602 emails, 432 of them are spam, and 2170 of them are ham.

Within the total of 48583 words, getting 8657 words occur in spam, and 45091 words occur in ham email.

```
1 print(train_X[0])
```

```
['letter', 'establish', 'convert', 'participate', 'start', 'customer', 'simple', 'inter']
```

```

1 # Build test data set
2 test_X = []
3 test_Y = []
4
5 test_path = 'lingspam_public/lemm_stop/part10'
6 test_mail_path = [(os.path.join(test_path, file_name), 'spmsg' in file_name) for file_name in os.listdir(test_path)]
7 for file_path, spam in test_mail_path:
8     with open(file_path) as f:
9         content = f.readlines()[2]
10        words = content.split()
11        filtered_words = list(set([w for w in words if (len(w) > 1 and re.match(word_filters, w))]))
12        # Add words to test set
13        test_X.append(filtered_words)
14        # Label spam as 1, ham as 0

```

```

15         if spam:
16             test_Y.append(1)
17         else:
18             test_Y.append(0)

1 # Helper parameter and function for IG
2 def Log(i):
3     return np.log(i) if i != 0 else 0
4
5
6 total_num = spam_num + ham_num
7 P_SPAM = spam_num / total_num
8 P_HAM = 1 - P_SPAM
9 HC = -P_SPAM * Log(P_SPAM) - P_HAM * Log(P_HAM)

1 # Generating IG for all words
2
3 IG = {}
4
5 for word in train_all_counter:
6     all_occur = train_all_counter[word]
7     spam_occur = train_spam_counter[word]
8     ham_occur = train_ham_counter[word]
9
10    # X=1, C = ham
11    P1h = (ham_occur / ham_num) * P_HAM
12    # X=0, C = ham
13    P0h = (1 - (ham_occur / ham_num)) * P_HAM
14    # X=1, C = spam
15    P1s = (spam_occur / spam_num) * P_SPAM
16    # X=0, C = spam
17    P0s = (1 - (spam_occur / spam_num)) * P_SPAM
18    # P(X=1)
19    P1 = all_occur / total_num
20    # P(X=0)
21    P0 = 1 - P1
22
23    # Note:  $H(C|X) = - \sum (P(X,C) \log(C|X))$ 
24    # But it is easier to do addition operation here :)
25    HCX = P1h*Log(P1h/P1) + P0h*Log(P0h/P0) + P1s*Log(P1s/P1) + P0s*Log(P0s/P0)
26    IG[word] = HC + HCX
27
28 top_words = [k for k, v in sorted(IG.items(), key=lambda item: -item[1])]
29 top_10 = top_words[:10]
30 top_100 = top_words[:100]
31 top_1000 = top_words[:1000]

1 print(top_10)

```

```
['language', 'remove', 'free', 'linguistic', 'university', 'money', 'click', 'market',
```

## ▼ Classifiers Implementation

### ▼ Generate Feature Matrix

```
1 def get_feature(N, dataset, term_frequency=False):
2     """Generate feature matrix based on top N words."""
3     # print(term_frequency)
4     top_n = {v:i for i,v in enumerate(top_words[:N])}
5     feature_matrix = np.zeros((len(dataset), N))
6     for j, content in enumerate(dataset):
7         for word in content:
8             if word in top_n:
9                 if term_frequency:
10                     feature_matrix[j, top_n[word]] += 1
11                 else:
12                     feature_matrix[j, top_n[word]] = 1
13
14     return feature_matrix
```

### ▼ Bernoulli NB Classifier

```
1 class BNB:
2     """Naive Bayes classifier for multivariate Bernoulli models."""
3     def __init__(self, alpha=1.0):
4         self.alpha = alpha # Actually, no need in this homework :)
5         self.spam_num = 0
6         self.ham_num = 0
7         self.P_X_spam = []
8         self.P_X_ham = []
9
10
11     def fit(self, train_x, train_y):
12         """Fit Naive Bayes classifier according to train_x, train_y."""
13         self.spam_num = sum(train_y)
14         self.ham_num = len(train_y) - sum(train_y)
15         rows, feature_num = train_x.shape
16
17         self.P_X_spam = np.zeros(feature_num)
18         self.P_X_ham = np.zeros(feature_num)
19         for i in range(feature_num):
```

```

20         x_spam = 0
21         x_ham = 0
22         for j in range(rows):
23             if train_x[j,i]:
24                 if train_y[j]:
25                     x_spam += 1
26                 else:
27                     x_ham += 1
28         self.P_X_spam[i] = (1 + x_spam) / (self.spam_num + 2)
29         self.P_X_ham[i] = (1 + x_ham) / (self.ham_num + 2)
30
31
32     def _row_predict(self, one_sample):
33         """Predict one record of the data set."""
34         prob = 1
35         for i, feature in enumerate(one_sample):
36             p0 = self.P_X_ham[i]
37             p1 = self.P_X_spam[i]
38             prob *= (p1 / p0) if feature == 1 else (1 - p1) / (1 - p0)
39         return prob * self.spam_num / self.ham_num
40
41
42     def predict(self, test_x):
43         """Perform classification on an array of test vectors test_x."""
44         pred = np.zeros((test_x.shape[0]))
45         for i in range(test_x.shape[0]):
46             pred[i] = int(self._row_predict(test_x[i]) > 1)
47         return pred
48
49
50     def score(self, test_x, test_y):
51         """Return the mean accuracy on the given test data and labels."""
52         y_pred = self.predict(test_x)
53         count = 0
54         for i in range(len(test_y)):
55             count += 1 if y_pred[i] == test_y[i] else 0
56         return count / len(test_y)
57

```

## ▼ Multinomial NB Classifier

```

1 class MNB:
2     """Naive Bayes classifier for multinomial models."""
3     def __init__(self, alpha=1.0):
4         self.alpha = alpha
5         self.spam_num = 0
6         self.ham_num = 0
7         self.P_X_spam = []
8         self.P_X_ham = []

```

```
9
10
11 def fit(self, train_x, train_y):
12     """Fit Naive Bayes classifier according to train_x, train_y."""
13     self.spam_num = sum(train_y)
14     self.ham_num = len(train_y) - sum(train_y)
15     rows, feature_num = train_x.shape
16
17     self.P_X_spam = np.zeros(feature_num)
18     self.P_X_ham = np.zeros(feature_num)
19     for i in range(feature_num):
20         x_spam = 0
21         x_ham = 0
22         for j in range(rows):
23             if train_x[j,i]:
24                 if train_y[j]:
25                     x_spam += 1
26                 else:
27                     x_ham += 1
28         self.P_X_spam[i] = (1 + x_spam) / (self.spam_num + 2)
29         self.P_X_ham[i] = (1 + x_ham) / (self.ham_num + 2)
30
31
32 def _row_predict(self, one_sample):
33     """Predict one record of the data set."""
34     prob = 1
35     for i, feature in enumerate(one_sample):
36         p0 = self.P_X_ham[i]
37         p1 = self.P_X_spam[i]
38         prob += (Log(p1) - Log(p0)) if feature == 1 else 0
39     return prob + Log(self.spam_num) - Log(self.ham_num)
40
41
42 def predict(self, test_x):
43     """Perform classification on an array of test vectors test_x."""
44     pred = np.zeros((test_x.shape[0]))
45     for i in range(test_x.shape[0]):
46         pred[i] = int(self._row_predict(test_x[i]) > 0)
47     return pred
48
49
50 def score(self, test_x, test_y):
51     """Return the mean accuracy on the given test data and labels."""
52     y_pred = self.predict(test_x)
53     count = 0
54     for i in range(len(test_y)):
55         count += 1 if y_pred[i] == test_y[i] else 0
56     return count / len(test_y)
```

## ▼ Classifiers Comparison

```

1 from sklearn.naive_bayes import MultinomialNB, BernoulliNB
2 from sklearn.metrics import classification_report
3
4 top_N = [10, 100, 1000]
5 result_accuracy = []
6
7 for N in top_N:
8     Xprint = get_feature(N, train_X)
9     # X label with term frequency
10    X_tf = get_feature(N, train_X, term_frequency=True)
11    Y = np.array(train_Y)
12
13    X_test = get_feature(N, test_X)
14    # X label with term frequency
15    X_test_tf = get_feature(N, test_X, term_frequency=True)
16    Y_test = np.array(test_Y)
17
18    model1 = BNB()
19    model2 = MNB()
20    model3 = MNB()
21
22    # train models
23    model1.fit(X, Y)
24    model2.fit(X, Y)
25    model3.fit(X_tf, Y)
26
27    # predict
28    y_1 = model1.predict(X_test)
29    y_2 = model2.predict(X_test)
30    y_3 = model3.predict(X_test_tf)
31
32    print(f'----- Using top {N} features -----')
33    print( '----- Bernoulli Naive Bayes -----')
34    print(classification_report(Y_test,y_1, target_names = ["ham", "spam"]))
35    print('--- Multinomial Naive Bayes with binary features ---')
36    print(classification_report(Y_test,y_2, target_names = ["ham", "spam"]))
37    print('---- Multinomial Naive Bayes with term frequency ---')
38    print(classification_report(Y_test,y_3, target_names = ["ham", "spam"]))
39
40    N_accuracy = [model1.score(X_test,Y_test),
41                  model2.score(X_test,Y_test),
42                  model3.score(X_test_tf,Y_test)]
43    result_accuracy.append(N_accuracy)

```

```

----- Using top 10 features -----
----- Bernoulli Naive Bayes -----
precision    recall  f1-score   support

ham          0.96      0.98      0.97      242

```

spam	0.89	0.82	0.85	49
accuracy			0.95	291
macro avg	0.93	0.90	0.91	291
weighted avg	0.95	0.95	0.95	291

--- Multinomial Naive Bayes with binary features ---

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

ham	0.98	0.95	0.97	242
spam	0.80	0.92	0.86	49
accuracy			0.95	291
macro avg	0.89	0.94	0.91	291
weighted avg	0.95	0.95	0.95	291

---- Multinomial Naive Bayes with term frequency ---

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

ham	0.98	0.95	0.97	242
spam	0.80	0.92	0.86	49
accuracy			0.95	291
macro avg	0.89	0.94	0.91	291
weighted avg	0.95	0.95	0.95	291

----- Using top 100 features -----

----- Bernoulli Naive Bayes -----

	precision	recall	f1-score	support
ham	0.94	1.00	0.97	242
spam	1.00	0.67	0.80	49
accuracy			0.95	291
macro avg	0.97	0.84	0.89	291
weighted avg	0.95	0.95	0.94	291

--- Multinomial Naive Bayes with binary features ---

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

ham	1.00	0.98	0.99	242
spam	0.89	1.00	0.94	49
accuracy			0.98	291
macro avg	0.95	0.99	0.96	291
weighted avg	0.98	0.98	0.98	291

---- Multinomial Naive Bayes with term frequency ---

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

ham	1.00	0.98	0.99	242
spam	0.89	1.00	0.94	49
accuracy			0.98	291

## ▼ Spam precision and spam recall



Params	Spam Precision	Spam Recall
Top 10 Words, Bernoulli Naive Bayes	0.89	0.82
Top 10 Words, Multinomial Naive Bayes with binary features	0.80	0.92
Top 10 Words, Multinomial Naive Bayes with term frequency	0.80	0.92
Top 100 Words, Bernoulli Naive Bayes	1.00	0.67
Top 100 Words, Multinomial Naive Bayes with binary features	0.89	1.00
Top 100 Words, Multinomial Naive Bayes with term frequency	0.89	1.00
Top 1000 Words, Bernoulli Naive Bayes	1.00	0.61
Top 1000 Words, Multinomial Naive Bayes with binary features	1.00	1.00
Top 1000 Words, Multinomial Naive Bayes with term frequency	1.00	1.00

## ▼ Compare with sklearn

```

1 result_accuracy_sk = []
2
3 for N in top_N:
4     model1_sk = BernoulliNB()
5     model2_sk = MultinomialNB()
6     model3_sk = MultinomialNB()
7
8     X = get_feature(N, train_X)
9     X_tf = get_feature(N, train_X, term_frequency=True)
10    Y = np.array(train_Y)
11
12    X_test = get_feature(N, test_X)
13    X_test_tf = get_feature(N, test_X, term_frequency=True)
14    Y_test = np.array(test_Y)
15
16    # train models
17    model1_sk.fit(X, Y)
18    model2_sk.fit(X, Y)
19    model3_sk.fit(X_tf, Y)
20
21    # predict
22    N_accuracy = [model1_sk.score(X_test, Y_test),
23                  model2_sk.score(X_test, Y_test),
24                  model3_sk.score(X_test_tf, Y_test)]
25    result_accuracy_sk.append(N_accuracy)
26
27
28 print(result_accuracy)
29 print(result_accuracy_sk)

```

```

[[0.9518900343642611, 0.9484536082474226, 0.9484536082474226], [0.9450171821305842, 0.9
[[0.9518900343642611, 0.9518900343642611, 0.9518900343642611], [0.9450171821305842, 0.9

```

## ▼ SVM based spam filter

```

1 from sklearn.model_selection import GridSearchCV
2 from sklearn.model_selection import train_test_split
3 from sklearn.svm import SVC
4 from sklearn.preprocessing import MinMaxScaler
5 from sklearn.base import clone
6
7
8 feature_sizes = [10, 100, 200]
9 modes = ["bf", "tf"]
10 gamma = [1, 0.1, 0.01, 0.001]
11 C = [0.1, 1, 10, 100]
12 degree = [2,3,10]
13 tuned_parameters = [{'kernel': ['rbf','sigmoid'], 'gamma': gamma, 'C': C},
14                      {'kernel': ['linear'], 'C': C},
15                      {'kernel': ['poly'], 'gamma': gamma, "degree" : degree, 'C':C}
16                      ]
17 now_best = {}
18 now_score = 0
19 best_clf = SVC()
20 for N in feature_sizes:
21     for mode in modes:
22         X = get_feature(N, train_X, term_frequency=(mode=='tf'))
23         X_test = get_feature(N, test_X, term_frequency=(mode=='tf'))
24         Y = np.array(train_Y)
25         if mode == 'tf':
26             ss = MinMaxScaler()
27             X = ss.fit_transform(X)
28             X_test = ss.transform(X_test)
29         clf = GridSearchCV(SVC(), tuned_parameters, cv = 5, scoring='recall_macro', verbose=0)
30         clf.fit(X, Y)
31
32         if clf.best_score_ > now_score:
33             now_score = clf.best_score_
34             now_best = {"size":N, "mode":mode, "score":clf.best_score_, "params":clf.best_estimator_.get_params()}
35             best_clf = clone(clf.best_estimator_)
36             print (now_best)
37
38 N = now_best["size"]
39 mode = now_best["mode"]
40 X = get_feature(N, train_X, term_frequency=(mode=='tf'))
41 X_test = get_feature(N, test_X, term_frequency=(mode=='tf'))
42 Y = np.array(train_Y)
43 Y_test = np.array(test_Y)
44
45 if mode == "tf":

```

```

46 ss = MinMaxScaler()
47 X = ss.fit_transform(X)
48 X_test = ss.transform(X_test)
49
50 best_clf.fit(X, Y)
51 pred_y = best_clf.predict(X_test)
52 print (now_best)
53 print (classification_report(Y_test, pred_y, target_names = ["ham", "spam"]))

```

```

{'size': 10, 'mode': 'bf', 'score': 0.9333773098447026, 'params': {'C': 1, 'gamma': 1,
{'size': 100, 'mode': 'bf', 'score': 0.9611211300362037, 'params': {'C': 1, 'gamma': 0.
{'size': 200, 'mode': 'bf', 'score': 0.9699466122688307, 'params': {'C': 10, 'gamma': 0
{'size': 200, 'mode': 'bf', 'score': 0.9699466122688307, 'params': {'C': 10, 'gamma': 0

```

	precision	recall	f1-score	support
ham	0.98	1.00	0.99	242
spam	0.98	0.88	0.92	49
accuracy			0.98	291
macro avg	0.98	0.94	0.96	291
weighted avg	0.98	0.98	0.98	291

## Methodology

I use `GridSearchCV` to help me find the best parameters for SVM. Also, we've see that higher Information Gain feature matrix can have a better result in prediction. So I compare the score between "BF" and "TF" and the score among "N={10, 100, 200}". The result shows that using **top 200** words, **binary features**, and {'C': 10, 'gamma': 0.1, 'kernel': 'rbf'} as the SVM parameter can get the best result.

### ▼ Adversarial Classification

### ▼ Attacker

```

1 class Attacker:
2     def __init__(self):
3         self.Lo = []
4         self.Lo0 = []
5         self.Lo1 = []
6
7
8     def cal_Lo(self, train_x, train_y):
9         rows, feature_num = train_x.shape

```

```

10     self.Lo = np.zeros(feature_num)
11     self.Lo0 = np.zeros(feature_num)
12     self.Lo1 = np.zeros(feature_num)
13
14     for i in range(feature_num):
15         x_spam = 0
16         x_ham = 0
17         for j in range(rows):
18             if train_x[j,i]:
19                 if train_y[j]:
20                     x_spam += 1
21                 else:
22                     x_ham += 1
23         P_X1_ham = x_ham / ham_num
24         P_X1_spam = x_spam / spam_num
25         P_X0_ham = 1 - P_X1_ham
26         P_X0_spam = 1 - P_X1_spam
27         self.Lo1[i] = Log(P_X1_spam/P_X1_ham)
28         self.Lo0[i] = Log(P_X0_spam/P_X0_ham)
29         self.Lo[i] = self.Lo1[i] - self.Lo0[i]
30
31
32     def _row_attack(self, one_sample):
33         sort_Lo = sorted(range(len(self.Lo)), key=lambda i: self.Lo[i])
34
35         cost = 0
36         sigmaLoX = 0
37         for i,v in enumerate(one_sample):
38             sigmaLoX += self.Lo1[i] if v==1 else self.Lo0[i]
39         if sigmaLoX < 0:
40             return 0, one_sample
41
42         for i in sort_Lo:
43             if(self.Lo[i] >= 0):
44                 return cost, one_sample
45             if one_sample[i]:
46                 continue
47             one_sample[i] = 1
48             sigmaLoX += self.Lo[i]
49             cost += 1
50             if (sigmaLoX < 0):
51                 return cost, one_sample
52
53
54     def attack(self, test_x, test_y):
55         costs = []
56         new_test_x = []
57         rows, feature_num = test_x.shape
58         for i in range(rows):
59             if test_y[i]:
60                 cost, changed = self._row_attack(test_x[i])

```

```

61         new_test_x.append(changed)
62         costs.append(cost)
63     else:
64         new_test_x.append(test_x[i])
65
66     return np.mean(costs), np.array(new_test_x)
67

```

## ▼ Defender

```

1 class Defender:
2     def __init__(self):
3         self.P_X1_ham = []
4         self.P_X1_spam = []
5         self.P_X0_ham = []
6         self.P_X0_spam = []
7         self.Lo = []
8         self.Lo0 = []
9         self.Lo1 = []
10        self.result = []
11
12    def fit(self, train_x, train_y):
13        rows, feature_num = train_x.shape
14        self.P_X1_ham = np.zeros(feature_num)
15        self.P_X1_spam = np.zeros(feature_num)
16        self.P_X0_ham = np.zeros(feature_num)
17        self.P_X0_spam = np.zeros(feature_num)
18        self.Lo = np.zeros(feature_num)
19        self.Lo0 = np.zeros(feature_num)
20        self.Lo1 = np.zeros(feature_num)
21
22        for i in range(feature_num):
23            x_spam = 0
24            x_ham = 0
25            for j in range(rows):
26                if train_x[j,i]:
27                    if train_y[j]:
28                        x_spam += 1
29                    else:
30                        x_ham += 1
31            self.P_X1_ham[i] = x_ham / ham_num
32            self.P_X1_spam[i] = x_spam / spam_num
33            self.P_X0_ham[i] = 1 - self.P_X1_ham[i]
34            self.P_X0_spam[i] = 1 - self.P_X1_spam[i]
35            self.Lo1[i] = Log(self.P_X1_spam[i] / self.P_X1_ham[i])
36            self.Lo0[i] = Log(self.P_X0_spam[i] / self.P_X0_ham[i])
37            self.Lo[i] = self.Lo1[i] - self.Lo0[i]
38

```

```

39
40 def _row_attack(self, one_sample):
41     sort_Lo = sorted(range(len(self.Lo)), key=lambda i: self.Lo[i])
42
43     cost = 0
44     sigmaLoX = 0
45     for i,v in enumerate(one_sample):
46         sigmaLoX += self.Lo1[i] if v==1 else self.Lo0[i]
47     if sigmaLoX < 0:
48         return 0, one_sample
49
50     for i in sort_Lo:
51         if(self.Lo[i] >= 0):
52             return cost, one_sample
53         if one_sample[i]:
54             continue
55         one_sample[i] = 1
56         sigmaLoX += self.Lo[i]
57         cost += 1
58         if (sigmaLoX < 0):
59             return cost, one_sample
60
61
62 def defence(self, model, attacked_x):
63     for sample in attacked_x:
64         if model._row_predict(sample):
65             self.result += [1]
66             continue
67         new_matrix = [[]]
68         for i in sample:
69             if i:
70                 new_row = [(x + [0]) for x in new_matrix] + [(x + [1]) for x in new_matrix]
71             else:
72                 new_matrix_f = [(x + [0]) for x in new_matrix]
73                 new_matrix = new_matrix_f
74         new_matrix = [np.asarray(x) for x in new_matrix]
75         realFroms = []
76         for origin in new_matrix:
77             _, attacked = self._row_attack(origin)
78             if np.array_equal(attacked, sample):
79                 if model._row_predict(origin):
80                     realFroms += [origin]
81         if len(realFroms):
82             P_spam = 0
83             P_ham = 0
84             for origin in realFroms:
85                 for i in origin:
86                     P_spam *= self.P_X1_spam[i] if i else self.P_X0_spam[i]
87                     P_ham *= self.P_X1_ham[i] if i else self.P_X0_ham[i]
88             P_spam += 1
89             P_ham += 1

```

```

90         self.result += [0 if P_spam < P_ham else 1]
91     else:
92         self.result += [0]
93
94     return np.array(self.result)

```

## ▼ Adversarial attack analysis

```

1 from sklearn.metrics import confusion_matrix
2
3 N = 10
4 X = get_feature(N, train_X)
5 Y = np.array(train_Y)
6 X_test = get_feature(N, test_X)
7 Y_test = np.array(test_Y)
8
9 # Use original model to predict Y_test
10 botnet = BNB()
11 botnet.fit(X, Y)
12 Y_pred_origin = botnet.predict(X_test)
13 print("Original classification report before attack")
14 print(classification_report(Y_test, Y_pred_origin, target_names = ["ham", "spam"]))
15 print("Confusion Matrix")
16 print(confusion_matrix(Y_test, Y_pred_origin))
17
18 # Assume that Attacker knows
19 # - original model's result: Y_pred_origin
20 # - original test words list: test_X
21 # - original test label: test_Y
22 # Using 'Add-Words' attack
23 hacker = Attacker()
24 hacker.cal_Lo(X, Y)
25 cost, X_test_attacked = hacker.attack(X_test, Y_test)
26 Y_pred_attacked = botnet.predict(X_test_attacked)
27 print("Original classification report after attack")
28 print(classification_report(Y_test, Y_pred_attacked, target_names = ["ham", "spam"]))
29 print("Confusion Matrix")
30 print(confusion_matrix(Y_test, Y_pred_attacked))
31 print(f"Average cost = {cost}")
32
33 # Defender
34 antispam = Defender()
35 antispam.fit(X, Y)
36 Y_pred_defence = antispam.defence(botnet, X_test_attacked)
37 print("Original classification report after defence")
38 print(classification_report(Y_test, Y_pred_defence, target_names = ["ham", "spam"]))
39 print("Confusion Matrix")

```

```
40 print(confusion_matrix(Y_test, Y_pred_defence))
41 print(f"Average cost = {cost}")
```

Original classification report before attack

	precision	recall	f1-score	support
ham	0.96	0.98	0.97	242
spam	0.89	0.82	0.85	49
accuracy			0.95	291
macro avg	0.93	0.90	0.91	291
weighted avg	0.95	0.95	0.95	291

Confusion Matrix

```
[[237  5]
 [ 9 40]]
```

Original classification report after attack

	precision	recall	f1-score	support
ham	0.83	0.98	0.90	242
spam	0.29	0.04	0.07	49
accuracy			0.82	291
macro avg	0.56	0.51	0.49	291
weighted avg	0.74	0.82	0.76	291

Confusion Matrix

```
[[237  5]
 [47  2]]
```

Average cost = 1.7959183673469388

Original classification report after defence


	precision	recall	f1-score	support
ham	0.00	0.00	0.00	242
spam	0.17	1.00	0.29	49
accuracy			0.17	291
macro avg	0.08	0.50	0.14	291
weighted avg	0.03	0.17	0.05	291

Confusion Matrix

```
[[ 0 242]
 [ 0  49]]
```

Average cost = 1.7959183673469388

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1272: UndefinedWarning:
  _warn_prf(average, modifier, msg_start, len(result))
```





---

 0s    completed at 5:51 PM

